


The AI Upload

— *Weekly Digest*

A curated read on artificial intelligence: what shipped, what changed, and what it means for the law.

08 SECTIONS • 21 STORIES • 12 Min READ

— THIS WEEK'S HEADLINES



Gemini 3.5 Live Translate Brings Natural Speech-to-Speech Across 70+ Languages

GOOGLE - THE KEYWORD


Hear translated speech in real time—voices keep their tone and pace across 70+ languages, with continuous output that skips awkward pauses.



Anthropic Debuts Claude Mythos 5 and Fable 5 With Tiered Access Controls

GEEKY GADGETS

A two-tier model launch where the most capable version stays locked behind government-backed access controls.



Apple Withholds AI Siri in EU, Citing Digital Markets Act Conflict

THE VERGE

Why 450 million Europeans won't get the new Siri—and why scholars doubt Apple's privacy defense holds up.

SECTION 01

Enterprise AI

OpenAI Files Confidential IPO, Joining SpaceX and Anthropic in AI Market Push

Read the article: **WIRED**

OpenAI confidentially filed paperwork for an initial public offering, the company announced Monday, joining a trio of AI players pursuing what could each be trillion-dollar debuts. Anthropic filed its confidential IPO paperwork on June 1, and Elon Musk's SpaceX filed publicly last month. OpenAI did not specify timing or how much it plans to raise, noting in a brief blog post that it expects the filing to leak.

The development carries significant legal complexity. OpenAI's nonprofit still owns roughly 25 percent of the company, more than \$200 billion, and retains power to block major decisions and fire executives. The company recently defeated Musk's lawsuit alleging it strayed from its nonprofit mission, but its structure remains under scrutiny by California and Delaware regulators. The Securities and Exchange Commission must now review the company's accounting and risk disclosures.

For legal professionals, the filing raises questions about how a public benefit corporation overseen by a nonprofit navigates SEC requirements while claiming it can weigh societal impacts over shareholder value. OpenAI remains unprofitable, with revenue between \$10 billion and \$20 billion last year offset by billions in losses. Public adv

Bezos-Backed Startup Flourish Seeks the Brain's 'Core Algorithm' for Efficient AI

Read the article: [Slashdot.org](#)

Jeff Bezos has invested in Flourish, a new neuro AI startup that has raised \$500 million and carries a reported \$2.5 billion valuation. Founded by former Amazon executive Rob Williams and neuroscientist Thomas Reardon, the company aims to study the brain's architecture to uncover its "core algorithm" and build AI that learns continuously while consuming far less energy than current large language models.

Reardon argues that today's frontier models, despite their power, are inefficient consumers of compute and data and remain static once trained. A person processes information on roughly 20 watts, he notes, while a single AI training chip uses more than 30 times that amount. Flourish's stated goal is a synthetic intelligence system that runs on 50 watts or less and adapts like a human mind, drawing on original wet lab neuroscience experiments conducted alongside AI research.

For legal professionals tracking AI's trajectory, the venture signals a possible shift away from data-hungry, energy-intensive models. Beyond Bezos, backers include Lux Capital, Google Ventures, and Catalio. The company plans to release interim models as products while pursuing its larger ambition.

As Costs Rise, Tech Firms Weigh Cheaper AI Models Over Frontier Giants

Read the article: [TechCrunch](#)

A core assumption underlying the AI industry, that the most powerful models are the ones worth using, may be starting to crack as costs mount and investor subsidies slow. With token prices rising, enterprise users are facing real cost pressure for the first time and beginning to weigh smaller, cheaper models against frontier ones.

The legal sector offers an early test case. Legal AI tool Harvey, working with inference platform Fireworks AI, cut its inference costs threefold without reducing quality by combining Claude Opus with Fireworks' GLM 5.1, reserving Opus for the most intensive tasks. Harvey co-founder Gabe Pereyra said quality remains paramount in legal work, but that the definition of quality is shifting toward using the most efficient model that still reaches the right answer.

Coinbase co-founder Brian Armstrong predicts that 80 percent of workloads will move to far cheaper models within 12 to 18 months. If true, the shift would reshape AI economics and deal a

financial blow to labs like OpenAI and Anthropic as they approach their IPOs.

Insiders Sound the Alarm as AI Costs Outpace Promised Savings

Read the article: [Axios](#)

A new phase of AI skepticism is emerging, and this time the doubters are insiders rather than outside critics. According to Axios, early corporate adopters are confronting the steep costs of treating AI as a universal productivity tool. Uber capped employee AI usage after exhausting its annual Claude Code budget in four months, with one executive calling the spending "harder to justify." Amazon shut down an internal token leaderboard after employees gamed it, and GitHub moved its Copilot assistant to usage-based billing to build a "sustainable" business.

The financial picture is sobering. A Bain survey of 951 large companies found AI savings falling well below projections even as most firms planned to increase spending, concluding that "the technology worked" but "the value didn't arrive." Even OpenAI CEO Sam Altman called the question of whether AI spending shows up in revenue "the most fair criticism" of the moment.

The debate carries market weight. Axios reports the Nasdaq fell 4.2% and the Philadelphia Semiconductor Index dropped 10.3% after Broadcom declined to raise its longer-term AI revenue outlook. For legal professionals tracking AI adoption costs and corporate risk, the full piece offers use

SECTION 02

Legal Intelligence

Big Law Overhauls Pricing, Hiring, and Services as AI Reshapes the Field

Read the article: [The Philadelphia Inquirer](#)

Major law firms are restructuring how they operate as artificial intelligence reshapes legal work, from pricing and hiring to the services they offer. The shift dates back several years. In 2021, Volkswagen approached global firm Freshfields to assess software compliance across more than 100 countries, a task that traditionally required local lawyers in each jurisdiction. Instead, Freshfields

Lab built a flexible AI platform allowing the carmaker to run real-time global risk assessments, a tool still in use today.

That early experiment is now widespread. Firms including Allen & Overy Shearman and Linklaters license custom AI products, while out-of-the-box giants Legora and Harvey have reached valuations of \$5.6 billion and \$11 billion, respectively. These tools enable rapid document review during due diligence and discovery and have introduced multitier pricing, threatening the billable hour as clients weigh speed and cost against perfection.

The risks remain real. A database tracking AI-introduced inaccuracies in legal filings has logged hundreds of instances, and Sullivan & Cromwell recently apologized for an AI-generated citation error. With routine tasks vanishing, firms anticipate slower junior hiring and pipeline concer

Musicians Union Sues UMG and WMG Over AI Training Settlements

Read the article: **Billboard**

The American Federation of Musicians has filed a lawsuit against Universal Music Group and Warner Music Group, alleging the major labels breached their collective bargaining agreement by licensing session musicians' recordings to AI companies Suno and Udio without compensation or credit. The complaint, filed June 5 and first obtained by Billboard, also claims the labels refused to disclose which recordings and whose work are being fed into the AI training sets.

The dispute stems from copyright settlements the majors reached at the end of 2025, allowing Suno and Udio to train new models on licensed works. AFM argues its collective bargaining agreement requires the labels to compensate members for "new uses" of their music and entitles the union to know which recordings are being used.

Both UMG and WMG responded that the union chose litigation amid ongoing collective bargaining negotiations, expressing disappointment and a desire to resolve matters at the table. The case spotlights an unresolved question for artists and their advocates: whether musicians will receive approval or compensation when their catalogs are used as AI training inputs, separate from the outputs the models generate. Compensation structures and the specific recordings involved remain unclear.

AI Products

Apple Pitches Privacy and Context to Developers in AI Comeback

Read the article: [The Register](#)

At its 2026 Worldwide Developers Conference, Apple presented a deliberately understated vision for integrating AI across its products, focusing on three areas: platform improvements, child safety enhancements, and Apple Intelligence. The company emphasized privacy, integration, and cost as its core advantages, with software engineering chief Craig Federighi calling privacy in AI "non-negotiable" and contrasting Apple's approach with providers that retain personal interactions by default.

For developers, the announcements carry practical weight. Apple is offering its Foundation Models framework, now based on Google's Gemini model family and newly multimodal, to run on-device or through Private Cloud Compute. Developers with fewer than two million first-time App Store downloads can use Apple Foundation Models in Private Cloud Compute with no cloud API costs. A new Core AI framework provides a memory-safe Swift API for running models entirely on-device with no token costs or server dependencies.

The update matters for legal professionals tracking AI and data privacy, since cloud-based models can capture sensitive information. Apple also expanded Xcode 27 to support Google's Gemini alongside Anthropic's Claude Agent and OpenAI's Codex, and r

Apple Unveils Gemini-Powered AI Overhaul at WWDC 2026 as Cook Steps Down

Read the article: [WIRED](#)

Apple's WWDC 2026 keynote centered on its long-anticipated artificial intelligence overhaul, the product of a new partnership with Google to use Gemini as a backbone for Apple Intelligence. Beginning this fall, AI features will be integrated across Mail, Messages, Safari, Photos, and a revamped Siri, with the assistant gaining broader contextual awareness, a new dedicated Siri AI app, camera-based visual search, and a more expressive voice. Apple says much of the processing occurs on-device, likely limiting the features to recent hardware, and emphasized that cloud requests are handled securely.

Notably for those tracking regulatory developments, Apple confirmed the new AI features will not launch in the European Union or China, citing unresolved regulatory hurdles tied to preserving its privacy protections. The company also introduced AI-driven photo editing tools that can reframe, extend, and alter the perspective of images, raising familiar questions about authenticity in digital media.

The event marked Tim Cook's final WWDC as CEO. Cook becomes executive chairman on September 1, with hardware engineering chief John Ternus taking over as chief executive. WIRED's full rundown covers iOS

Google Brings Latest Gemini Models to Apple Developers via Foundation Models Framework

Read the article: [Google - The Keyword](#)

Google announced at Apple's Worldwide Developers Conference that Apple developers can now call cloud-hosted Gemini models through the Foundation Models framework and access Gemini directly in Xcode. Apple is opening its Foundation Models framework to third-party cloud model providers starting with iOS 27, macOS 27, iPadOS 27, visionOS 27 and watchOS 27, and Google has made Gemini available through the Firebase Apple SDK.

The integration places Apple's on-device model and cloud-hosted Gemini models behind a shared API surface, allowing developers to swap between local and cloud inference. Switching to Gemini requires only a small code change. The setup runs on Firebase AI Logic, removing the need for a separate backend server, while Firebase App Check protects service APIs from abuse.

The Xcode integration supports multi-step agentic tasks such as code review and bug fixing. Individual developers can obtain a self-serve API key through Google AI Studio, with free and paid tiers, while enterprise developers can use the Gemini Enterprise Agent Platform for dedicated corporate quotas and data privacy controls. The Foundation Models integration is available as a preview release.

Security

Anthropic's Mythos Converts Software Patches Into Working Exploits Within Minutes

Read the article: [Axios](#)

Anthropic's new research suggests that advanced AI models can rapidly convert publicly disclosed software vulnerabilities into working exploits, a capability that could shrink the so-called "patch gap" between when a flaw is disclosed and when systems are widely updated. The company's Mythos Preview turned newly disclosed bugs into functional exploits in hours rather than weeks, according to findings shared first with Axios.

Anthropic's frontier red team tested Mythos against vulnerabilities disclosed in January and February in Mozilla Firefox and the Microsoft Windows kernel, focusing on bugs disclosed after the models' knowledge cutoff dates. Within 31 minutes, Mythos generated its first proof-of-concept exploit for a Windows kernel vulnerability, and it caused a "blue screen of death" in 18 of 21 kernel bugs tested. On Firefox, it built eight working code-execution exploits across 18 security patches. Anthropic estimated the Windows privilege-escalation exploits cost roughly \$2,000 each in API credits.

The concern extends beyond Mythos, as some open-source models and OpenAI's GPT-5.5-Cyber are reportedly performing at

Self-Replicating AI Worm Runs Entirely on Local Open-Weight Models

Read the article: [The Hacker News](#)

University of Toronto researchers have built an AI-powered computer worm that can break into networks on its own, no human guidance required. In lab tests on a 33-machine network, the worm found vulnerabilities, gained access to systems, and spread itself to 62 percent of the network within seven days.

Unlike traditional malware, which relies on a fixed set of pre-programmed attack methods, this worm thinks on the fly. It reads the latest public security advisories and adapts its approach to whatever it finds on each new target. That means patching a single known flaw is not enough to stop it.

What makes it especially difficult to contain is that it runs on freely available AI models, with no connection to commercial AI providers. There is no account to suspend, no API key to cut off. The researchers have not released the code publicly, but they argue the findings should push organizations to rethink how quickly they respond to newly disclosed vulnerabilities, because the window between a flaw being published and it being exploited may now be measured in hours.

ChatGPT's New Lockdown Mode Restricts Tools to Curb Data Exfiltration

Read the article: [The Hacker News](#)

OpenAI has begun rolling out a Lockdown Mode for eligible personal ChatGPT accounts, an optional security setting aimed at reducing the risk of data exfiltration from prompt injection attacks. The feature targets individuals and organizations that handle sensitive data and need stricter protection. It is available to logged-in users on Free, Go, Plus, and Pro tiers, as well as self-serve ChatGPT Business plans.

Rather than preventing prompt injections outright, Lockdown Mode limits the pathways through which data could leave the system by restricting outbound network requests. Activating it disables or limits several features, including live web browsing, image support, deep research, agent mode, Canvas networking, and file downloads. The mode cannot run alongside Developer Mode. OpenAI cautioned that the setting does not guarantee data exfiltration cannot occur, noting that risks may persist through enabled apps, unforeseen capability combinations, or newly discovered techniques.

For legal professionals managing confidential client information, the trade-off between functionality and security is worth understanding. OpenAI also introduced an account management feature allowing users to review active sessions and log out of individual or all devices if unauthorized activity is suspected

SECTION 05

Policy

White House and Congress Revive Push to Override State AI Laws

Read the article: [Axios](#)

The White House is negotiating a deal with congressional lawmakers to preempt certain state AI laws in exchange for backing key tech policy priorities, according to Axios. Sen. Marsha Blackburn (R-Tenn.) is leading the talks, which would pair federal override of state AI regulations with the Kids Online Safety Act, the NO FAKES Act, and age verification requirements. A Blackburn spokesperson stressed the package is not a "blanket preemption of all laws regulating AI or child safety."

The development matters because states are increasingly passing stronger AI laws, and the Trump administration faces pressure to act. Blackburn's support, which the White House previously lacked, would be central to passage. Her involvement also suggests that a competing bipartisan bill from Reps. Jay Obernolte (R-Calif.) and Lori Trahan (D-Mass.), which would preempt state AI laws for three years and establish the Center for AI Standards and Innovation, is not the likely vehicle.

The renewed effort follows Trump's recent AI and cyber executive order, which includes voluntary pre-deployment testing of frontier models. Axios reports the White House is also meeting with AI companies this week to discuss benchm

UK Commits \$1.47 Billion to AI Supercomputer, Cutting Reliance on US Tech

Read the article: **WIRED**

The UK government has unveiled a \$1.47 billion plan to reduce its dependence on foreign-made artificial intelligence hardware. Announced Monday, the measures include more than \$1 billion for a national AI supercomputer, stocked with \$530 million worth of hardware. Of that, \$200 million will go toward specialist inference chips for processing AI tasks. The government will give procurement priority to British firms, naming startups Olix and Fractile as potential beneficiaries. Researchers and startups are expected to access the supercomputer beginning in 2030.

The plan forms part of a broader push to minimize reliance on foreign powers for AI products and services, an effort made more urgent by tensions between the US and its European counterparts. The European Union outlined a similar tech sovereignty proposal last week. UK technology secretary Liz Kendall framed AI sovereignty as a matter of "reducing overdependencies and increasing resilience."

The hardware plan joins earlier UK initiatives, including AI growth zones launched last November and the \$675 million SovAI venture fund launched in April. Observers say the government's

willingness to back domestic firms with hard contracts could anchor promising chip startups in the country and carve out a strategically important niche.

Trump's Executive Orders Signal Push for AI Growth Over Safety Constraints

Read the article: [The Guardian - Technology](#)

President Trump's recent actions signal a clear commitment to accelerating AI development rather than constraining it for safety. According to the article, Trump issued two executive orders last week, one seeking a government review of AI models 30 days before release without mandating it, and another directing the defense department to accelerate AI adoption for national cybersecurity. He had earlier postponed signing the review order, citing concerns about maintaining the US lead over China.

The administration is also reportedly discussing taking financial stakes in leading US AI companies, with OpenAI's Sam Altman said to have participated in talks with White House officials. Trump described the concept as "almost a partnership with the American public."

Meanwhile, Anthropic advocated for a possible "temporary pause" on advancing AI capabilities, even as it confidentially filed for an initial public offering. Valued around \$965bn, the company now surpasses rival OpenAI. The article raises pointed questions about who such a freeze would apply to and whether self-described ethical labs would benefit. For legal professionals tracking AI regulation and corporate governance, these developments warrant attention.

SECTION 06

Responsible AI

Stanford Study Finds AI Hiring Tools Disadvantage Black and Asian Applicants

Read the article: [The Root](#)

Algorithmic hiring tools, now used by more than 90 percent of companies worldwide for initial candidate screening, may be amplifying racial discrimination at a scale that individual rejections never could. That is the central finding of a new study from Stanford University's Human-Centered

Artificial Intelligence institute, which the researchers describe as the first large-scale analysis of hiring algorithms operating in real-world conditions.

The team examined more than 4 million job applications from 3.4 million people across 156 employers using the same AI hiring vendor. They found the tools effectively locked out nearly 26 percent of Black candidates and also produced negative effects for Asian applicants. Had these groups been recommended at the same rate as the most-favored group, typically white applicants, 40,000 more applications would have advanced. The study also identified offsetting patterns, with Black candidates less likely to be recommended in finance but more likely in warehouse roles, which can make discrimination appear absent.

The researchers tied these outcomes to market concentration and warned that widespread use, high-stakes consequences, and limited transparency rarely belong together. They call for independent research to support evidence-based AI policy and accountability.

— SECTION 07

AI Sustainability

China Launches World's First Wind-Powered Underwater Datacentre off Shanghai

Read the article: [**The Guardian - Technology**](#)

China has launched the world's first wind-powered underwater datacentre off the coast of Shanghai, advancing the country's efforts to address energy challenges driven by its artificial intelligence expansion. The Shanghai Lingang undersea project, which began operations in May, has a capacity of 24 megawatts and is a joint effort between HiCloud Technology and the state-owned China Communications Construction.

Submerged 10 metres below the surface more than 6 miles offshore, the facility draws power from a nearby offshore windfarm. The Chinese government reports it reduces power consumption by over one-fifth compared with land-based datacentres, partly because seawater provides natural cooling. Traditional facilities can spend 25% to 40% of their electricity on chilling servers and have faced scrutiny over their freshwater use. The United Nations University recently warned that datacentre water consumption could reach 9.3tn litres by 2030.

The project, which received 1.6bn yuan in investment, builds on HiCloud's 2023 commercial underwater datacentre in Hainan and follows an earlier Microsoft pilot in Scotland. Experts noted potential risks to marine ecosystems, including disturbed sediment and warmer seawater

SECTION 08

Creative AI

Suno's \$5.4 Billion Valuation Signals AI Music Boom Amid Unresolved Copyright Disputes

Read the article: **Fortune**

AI music generator Suno announced a new \$400 million funding round at a \$5.4 billion valuation, signaling continued investor confidence in AI-generated music. The platform allows users to create songs from simple text prompts in seconds, and according to fundraising materials reported by Billboard, users were generating more than 7 million songs per day at the time of the financing.

For the legal community, the valuation arrives against a backdrop of unresolved copyright litigation. Suno and rival Udio have acknowledged training their models on copyrighted recordings, arguing the practice qualifies as fair use, a position copyright holders dispute. Universal Music Group, Sony Music, and Germany's GEMA continue to pursue legal action, while Warner Music Group reached a licensing agreement last year. When the labels first sued Suno in 2024, they alleged training on roughly 560 copyrighted works. Last month they sought to amend their complaint to claim more than 61,000 additional songs were used without permission. Both Suno and Udio have asked courts to keep their training dataset sizes confidential.

These questions about lawful training and demand remain open even as investors look past the uncertainty.



This newsletter contains AI-generated content that may contain inaccuracies. For research or citation purposes, please read and cite the original source article, not this newsletter.