

The AI Upload

— *Weekly Digest*

A curated read on artificial intelligence: what shipped, what changed, and what it means for the law.

09 SECTIONS · 23 STORIES · 14 Min READ

— THIS WEEK'S HEADLINES



Senate Advances NO FAKES Act to Curb AI Deepfakes, Civil Liberties Concerns Persist

CYBERSCOOP


Lawmakers want to give you rights over your AI likeness—but free speech advocates warn of unintended censorship.



Anthropic's Mythos Export Ban Grows More Tangled Amid Security Disputes

THE REGISTER

Ninety minutes to comply, no explanation given—and security experts say the data behind the ban points the other way.



Google's Interactions API Hits General Availability as Primary Gemini Interface

GOOGLE - THE KEYWORD

Managed agents, remote sandboxes, and background tasks—what the new default Gemini interface means for developers building agentic workflows.

SECTION 01

Enterprise AI

Google Loses Key DeepMind Researchers Amid Escalating AI Talent War

Read the article: [Axios](#)

Google DeepMind suffered two significant researcher departures within a week. Noam Shazeer, co-author of the 2017 "Attention Is All You Need" paper that introduced the Transformer architecture, announced he is leaving for OpenAI after Google paid more than \$2 billion to acquire him and part of his Character.ai team. Days later, John Jumper, who shared the 2024 Nobel Prize in Chemistry for his work on AlphaFold, announced his departure for Anthropic.

The moves reflect broader turbulence across major AI labs, with Nvidia acquiring the team behind Essential AI and Barret Zoph departing OpenAI for a second time under unexplained circumstances. According to the article, top researchers bring more than technical skill: they offer judgment about which ideas to pursue, experience managing large-scale experiments, and the ability to attract additional talent.

For legal professionals tracking AI governance and industry concentration, the competitive dynamics here carry real stakes. Anthropic and OpenAI reportedly hold a financial advantage through anticipated IPOs, while researchers weighing where to work factor in each organization's leadership and its approach to responsible development.

Alphabet Stock Drops Sharply Amid Key AI Departures and Competitive Pressures

Read the article: [CNBC](#)

Alphabet shares fell roughly 5% on June 22, marking the company's worst single-day stock decline in over a year, as a combination of talent losses, competitive pressure, and service outages rattled investors. The slide was triggered in part by the departure of Noam Shazeer, Google's vice president of engineering and co-lead of its Gemini AI models, who announced he was joining OpenAI. Days later, John Jumper, a DeepMind vice president and Nobel Prize winner best known for co-creating AlphaFold, announced he was leaving for Anthropic after nine years at the company.

The sell-off also followed a Wall Street Journal interview in which Microsoft CEO Satya Nadella described the AI market as commoditized and urged less reliance on major AI providers. That framing put added scrutiny on Alphabet's aggressive spending, which has included raising \$141 billion in debt and equity since October. Gmail and YouTube outages on the same day compounded an already difficult news cycle for the company.

— SECTION 02

Legal Intelligence

Inside the 21 Million Tracks Used to Train Generative AI Music Models

Read the article: [The Atlantic](#)

Four large music datasets circulating within the AI development community contain a combined total of more than 21 million tracks, according to an Atlantic investigation published in June 2026. The datasets include recordings by artists ranging from Taylor Swift and the Beatles to Miles Davis and Vijay Iyer, and span genres from pop and jazz to classical. The 12-million-track dataset alone would take 91 years to listen to in full. Google and Stability AI have each been documented using at least one of the datasets for model training.

The investigation also found that AI music generators, including Suno and Udio, have produced tracks that closely resemble copyrighted songs by Michael Jackson, Chuck Berry, and Ed Sheeran, among others. Suno stated in a 2024 court filing that it trained its models on "essentially all music files of reasonable quality" available on the internet. At least 12 copyright infringement lawsuits have been filed against AI companies by musicians and major record labels, though no rulings have been issued. Sony recently identified 135,000 AI-generated tracks attributed to its artists on

streaming platforms. The article raises significant questions for legal professionals tracking how fair use doctrine applies to AI training at scale.

UK's First AI Law Firm Wins Debut Court Case for £400 Fee

Read the article: [Slashdot.org](#)

Garfield AI, described as the UK's first regulator-approved AI law firm, secured a court victory at Wandsworth County Court on May 14, 2026, helping a freelancer recover 7,000 pounds in unpaid fees from a hospitality business. The claimant, Tamires Camal Taquidir, paid approximately 400 pounds in Garfield AI fees for the representation, while the opposing party retained both a solicitor and a barrister.

Garfield AI handled all pre-trial legal work, including generating pre-action correspondence, preparing and issuing court proceedings, managing document production, and drafting witness statements and trial bundles. A junior barrister was instructed shortly before trial to appear in court. The defendant's counterclaim was dismissed, and the court awarded the full 7,000 pounds to the claimant.

The case is notable for demonstrating a cost-access gap: the claimant spent roughly 400 pounds to recover 7,000 pounds, while the opposing side employed traditional legal counsel at presumably far greater expense. For legal professionals, the outcome raises practical questions about how AI-assisted legal services may affect access to justice and competitive dynamics in routine civil litigation.

Legora Hits \$5.6 Billion Valuation as Legal AI Startup Expands Globally

Read the article: [Fortune](#)

Swedish legal AI startup Legora has reached a \$5.6 billion valuation following a \$550 million Series D fundraise, and now serves 1,000 law firms and in-house legal teams across 50 markets. Founded three years ago by self-described "engineers and hackers," the company crossed the \$100 million annual recurring revenue threshold this year. Backers include NVentures, Nvidia's corporate venture arm, and law firm Bird & Bird, which also advises on product development. Linklaters and Cleary Gottlieb are among its clients.

The company drew widespread attention through an advertising campaign featuring actor Jude Law, who insisted on bringing the cinematographer from *Oppenheimer* and a *Saturday Night Live*

scriptwriter. CEO Max Junestrand credited the campaign's origin to a wine-fueled internal brainstorm and describes persuading Law's representatives as an exercise in "nagging as negotiation."

Bird & Bird CEO Christian Bartsch offered a frank assessment of why legal is ripe for AI: fragmented systems, middleware patches, and labor-intensive workflows sitting beneath high-value strategic work. His warning to the profession is pointed: law firms that use AI well will replace those that do not, with the next generation of lawyers operating as human-machine hybrids.

NYT CEO Emphasizes High Stakes in Copyright Lawsuit Against OpenAI

Read the article: [Crypto Briefing](#)

The New York Times CEO Meredith Kopit Levien has been speaking publicly about the high stakes of the newspaper's copyright lawsuit against OpenAI and Microsoft, originally filed in December 2023 in the Southern District of New York. The suit alleges that millions of Times articles were used to train large language models, including those powering ChatGPT, without authorization or compensation. Initial damages are estimated at a minimum of \$2.25 billion, with some observers suggesting the total could rise significantly higher.

Levien has framed the litigation as a test case for the entire AI industry, arguing the outcome will determine whether AI companies must compensate content creators whose work trained their models. The Times has also filed a separate suit against Perplexity AI and is reportedly involved in additional OpenAI investigations. Notably, the newspaper has simultaneously pursued licensing deals with other tech companies, including a reported agreement with Amazon, signaling that licensing revenue may be the Times' ultimate objective rather than halting AI development altogether.

SECTION 03

AI Products

Nvidia Unveils Agentic AI Stack for Scientific Supercomputing at ISC 2026

Read the article: [The Register](#)

Nvidia used the ISC High Performance 2026 conference in Hamburg to announce a new scientific computing stack built around agentic AI, positioning autonomous systems as the future of large-scale research computing. The company introduced three new software tools: ALCHEMI, a toolkit for chemical and materials discovery; DAQIRI, which connects scientific instruments to real-time AI inference and is already being tested at CERN's ATLAS experiment to capture more collision data; and cuPhoton, designed to process telescope and camera data at dramatically accelerated speeds. Nvidia claims cuPhoton achieved a 15,000-fold improvement in image loading and up to 8,000 times faster signal processing in tests simulating data from the Rubin Observatory.

The hardware centerpiece is the forthcoming Vera Rubin NVL rack, expected in Q4 2026, offering up to 144 GPUs per rack and 5 petaFLOPS of FP64 performance. The Los Alamos National Laboratory's Mission and Vision supercomputers will be among the first deployments, and Nvidia describes them as the world's first agentic AI supercomputers. For legal professionals tracking AI governance and research infrastructure, the expansion of autonomous AI systems into federally funded scientific institutions raises questions about oversight, accountability, and the evolving role of human researchers in AI-assisted discovery.

Agentic AI Loops Promise Continuous Automation at Open-Ended Cost

Read the article: [TechCrunch](#)

Boris Cherny, creator of Claude Code, told attendees at Meta's @Scale conference that agentic loops represent a shift as significant as the move from hand-written code to AI-generated code. Rather than a user directing a single agent toward a discrete task, loops set multiple agents running continuously in the background, with one agent prompting another. Cherny described running such loops in his own work, where separate agents perpetually scan his codebase for architectural improvements and duplicated abstractions, submitting pull requests like human collaborators.

The concept builds on familiar recursive logic from computer science but adds a non-deterministic twist: a subagent, rather than a fixed condition, decides when to stop. Popular implementations like the "Ralph Loop" keep models on track during long runs by repeatedly asking whether the goal has been achieved. Researchers at OpenAI have noted that sufficient compute can solve nearly any problem, and loops operationalize that idea by simply never stopping.

The catch is cost. Agentic loops consume tokens far faster than conventional chatbot interactions, with no natural ceiling on spending. For legal professionals and organizations evaluating AI

workflows, that tradeoff between continuous autonomous improvement and open-ended computational expense may be the central question worth watching.

SECTION 04

Security

Five Eyes Warn Frontier AI Will Transform Cybersecurity Within Months, Not Years

Read the article: [CyberScoop](#)

The Five Eyes intelligence alliance, including agencies from the United States, Canada, the United Kingdom, Australia, and New Zealand, has issued a joint warning that advanced AI models capable of executing sophisticated cyberattacks will be publicly available within months. Signed by NSA Cybersecurity Directorate Director David Imbordino and acting CISA Director Nick Andersen, the statement warns that frontier models are expected to "fundamentally transform both offensive and defensive cyber capabilities" on a timeline of months, not years.

The agencies identify legacy systems, weak identity controls, slow patching, and lack of incident planning as prime targets that AI will be especially effective at exploiting. They note that AI models with serious cybersecurity applications already circulate through commercial, open-source, and black-market channels. Representative Andrew Garbarino of the House Homeland Security Committee added that China may be weeks away from achieving comparable frontier AI capabilities.

Despite the urgency, the agencies' core recommendations remain familiar: patch systems, limit unnecessary connectivity, and treat cybersecurity as a strategic priority rather than a compliance exercise. Defensive programs such as Anthropic's Project Glasswing and OpenAI's Trusted Access for Cyber Program aim to give organizations a head start in identifying vulnerabilities before adversaries can exploit them at scale.

OpenAI Launches Open-Source Bug-Fixing Program Amid Anthropic Export Controls

Read the article: [WIRED](#)

OpenAI announced a cluster of cybersecurity initiatives on Monday, including an updated version of its restricted GPT-5.5-Cyber model, expanded government partnerships under its Trusted Access for Cyber program, and a plug-in release of its Codex Security scanner. The headline initiative is Patch the Planet, a collaboration with security firm Trail of Bits and vulnerability management companies HackerOne and Calif., aimed at providing free security consulting to open-source software maintainers. The effort has already engaged more than 30 projects, uncovering hundreds of bugs and producing dozens of patches in its first week.

The program addresses a specific pressure point: AI-powered vulnerability discovery tools are generating large volumes of bug reports faster than volunteer maintainers can process them, crowding out attention from genuinely critical flaws. Patch the Planet offers individualized support to help maintainers triage reports, build patches, and integrate AI security tooling into their workflows. Participants receive six months of free ChatGPT Pro and Codex Security access.

The announcements arrive as competitor Anthropic faces U.S. export controls on its Mythos 5 and Fable 5 models over cybersecurity capability concerns, and as the Five Eyes intelligence alliance issued a joint warning that frontier AI models could transform offensive and defensive cyber capabilities within months.

SECTION 05

Policy

Trump Calls Anthropic a National Security Threat, Then Signals Thaw

Read the article: [Axios](#)

President Trump revealed in an exclusive interview with Axios that he considered Anthropic a national security threat as recently as a week before the interview, though he indicated the relationship has since improved following his interactions with CEO Dario Amodei at the G7 summit. The administration had imposed sweeping export controls restricting foreign access to Anthropic's most advanced models and the Pentagon designated the company a supply chain risk, treatment the article notes is typically reserved for foreign adversaries. The concerns originated from an Amazon report detailing a security vulnerability, which administration officials brought to Anthropic leadership but felt was dismissed.

Trump stopped short of ruling out invoking the Defense Production Act, while noting he did not expect to use it. The two sides are now reportedly collaborating on standards to evaluate AI jailbreaks. Anthropic issued a statement expressing gratitude for the administration's partnership and reaffirming its commitment to U.S. AI leadership. The episode highlights how quickly domestic AI companies can find themselves subject to national security scrutiny, with significant implications for AI governance, export policy, and the broader regulatory environment facing the industry.

NYC House Primary Becomes Key Battleground in the 'AI Civil War'

Read the article: [**The Guardian - Technology**](#)

AI-focused Super PACs have poured nearly \$16 million into a single New York City congressional primary, making Tuesday's Democratic contest in NY-12 the most concentrated battleground in what operatives are calling the "AI civil war." At the center is Democratic assemblymember Alex Bores, who sponsored state legislation requiring major AI developers to publish public safety plans. That move drew an \$8.2 million attack campaign from Think Big, an affiliate of Leading the Future, a Super PAC funded by venture capitalists Marc Andreessen and Ben Horowitz and OpenAI co-founder Greg Brockman. The group favors federal AI regulation over state-level frameworks.

The spending triggered a counter-offensive from rival Super PACs, including Public First, whose disclosed backers include Anthropic, which contributed \$20 million. Two-thirds of House Democratic AI policy leadership now has Public First funding behind their campaigns. Both camps have also spent heavily in races tied to rural datacenter expansion across Utah, Texas, Ohio, Georgia, and Kentucky.

The dynamic has broad implications for legal professionals tracking AI governance. With AI polling as politically unpopular across party lines and the first generation of federal AI legislation taking shape, Tuesday's results may signal how much industry money can shape the regulatory landscape.

AI CEOs Seated as Peers to World Leaders at G7 Summit

Read the article: [**Axios**](#)

At this week's G7 summit in Évian-les-Bains, France, the CEOs of OpenAI, Anthropic, Google DeepMind, and other major AI companies sat alongside heads of state as functional peers, a seating arrangement that Axios describes as capturing "a once-unimaginable geopolitical ordering." Sam Altman held bilateral meetings with multiple world leaders, while Dario Amodei, Altman, and

Meta's Alexandr Wang each posed with French President Macron in chairs typically reserved for heads of government.

Inside the closed working lunch, all three AI executives urged democratic nations to present a unified front on AI development. Altman called for an international standards body to establish testing protocols and governance frameworks, while cautioning that "no single lab should be making the decisions." Amodei warned democracies to "resist the temptation to splinter," and Google DeepMind CEO Demis Hassabis framed the moment as the "foothills of the singularity," advocating for a U.S.-led standards body operating in coordination with democratic allies.

For legal professionals tracking AI governance, the summit signals that the regulatory and geopolitical frameworks shaping AI development are increasingly being negotiated in real time between governments and private companies at the highest levels of international diplomacy.

— SECTION 06

Responsible AI

Rights Groups Slam Home Office AI Age Estimation Tool for Asylum-Seeking Children

Read the article: [The Register](#)

More than 60 human rights organizations, including Amnesty International, Human Rights Watch, and the Electronic Frontier Foundation, have signed an open letter urging the UK Home Office to abandon its plans to deploy AI-powered facial age estimation (FAE) technology on asylum-seeking children. The technology, slated for rollout from 2027, would help immigration officers determine whether individuals claiming to be minors are likely over or under 18. Ministers maintain it will support rather than replace human decision-making, but the coalition argues the system is biased, inaccurate, and potentially unlawful.

The letter highlights that the Home Office's own guidance acknowledges the technology performs differently across ethnicities and skin tones, raising discrimination concerns given that asylum-seeking children are predominantly people of color. Compounding the problem, the best-performing systems carry an error margin of roughly 2.5 years precisely at the 16-to-18 age boundary the government wants assessed. The groups also note that trauma, malnutrition, and

other hardships common among asylum-seeking children can make them appear older, further skewing results.

The organizations have given the Home Office 21 days to answer detailed questions about training data, testing methodologies, and safeguards. No Equality Impact Assessment or Data Protection Impact Assessment has been made public.

Tracking 2026's Major Tech Layoffs Where Employers Blamed AI

Read the article: [TechCrunch](#)

Oracle's June 22 annual filing revealed the company reduced its workforce by 21,000 employees over the past 12 months, a 13% decline, citing AI adoption as a contributing factor. The disclosure adds fresh specificity to a broader pattern TechCrunch has been tracking: major tech companies reporting strong revenue growth while simultaneously reducing headcount and attributing cuts, at least in part, to artificial intelligence.

The running list TechCrunch compiled spans more than a dozen companies since January 2026, including Amazon (16,000 corporate jobs), Meta (8,000 employees), PayPal (projected 4,500+ over two to three years), and Block (4,000 jobs, nearly half its workforce). Cloudflare's CEO described the majority of those laid off as "measurers," while Atlassian's CEO acknowledged AI "changes the number of roles required in certain areas." Outplacement firm Challenger, Gray & Christmas reported that tech layoffs hit their highest single month in years in May 2026, with AI as the most-cited reason.

Legal professionals should take note: the affected roles span not only engineering but also legal, compliance, internal auditing, HR, and customer support functions, suggesting AI-driven restructuring is reaching into areas directly relevant to the legal industry.

Doctorow's "Reverse Centaur" Examines AI's True Cost to Workers and Markets

Read the article: [The Guardian - Technology](#)

Cory Doctorow's latest book, "The Reverse Centaur's Guide to Life After AI," arrives amid growing public hostility toward artificial intelligence, which a majority of Americans now believe will negatively affect jobs, creativity, and human relationships. Doctorow's central metaphor distinguishes between the "centaur," a worker empowered by AI tools, and the "reverse centaur,"

one whose autonomy is diminished by machine demands. He argues that while the technology itself permits the former, current business models consistently produce the latter.

Doctorow's core critique targets the AI industry's revenue model rather than the technology itself. He contends that colossal sector valuations rest on the promise of replacing human labor, and that everyday users of AI products function less as customers than as unwitting participants in a hype cycle aimed at investors. He also warns that a collapse of this apparent bubble, given that seven major tech companies represent one-third of U.S. stock market value, could produce an economic shock comparable to 2008 or 2020. For legal professionals tracking AI's trajectory in labor, intellectual property, and corporate governance, the book offers a pointed structural framework worth examining.

SECTION 07

AI Sustainability

Microsoft and Chevron Plan Major Gas-Powered Data Center in West Texas

Read the article: [TechCrunch](#)

Microsoft and Chevron announced plans Monday for a 2.67-gigawatt natural gas power plant in West Texas, described by Chevron as among the largest co-located natural gas power and data center developments in the United States. Under a 20-year power purchase agreement, the plant will supply dedicated electricity to a Microsoft-operated data center supporting its AI and cloud services. Power generation will rely primarily on two GE Vernova turbines, with additional capacity from Solar Turbines, a Caterpillar subsidiary.

The project, known as Project Kilby, carries significant environmental implications. The Environmental Integrity Project estimates it could release more than 13 million tons of carbon dioxide, 3,200 tons of criteria air pollutants, and 278,000 pounds of hazardous air pollutants. That sits uneasily alongside Microsoft's stated pledge to eliminate its carbon emissions by 2030. Legal professionals tracking corporate climate commitments, ESG disclosure obligations, and energy regulatory frameworks will want to watch how this development is characterized in Microsoft's future sustainability reporting.

Nvidia's Cooling System Cuts Data Center Water Use, But Not AI's Full Footprint

Read the article: [TechCrunch](#)

Nvidia recently announced a warm-water cooling system it claims can eliminate "pretty much all water usage" inside data centers, using a closed-loop coolant that circulates at 45°C through server racks without requiring evaporative cooling or fans. The company's chief sustainability officer told Axios that "the water consumption challenge for data centers is largely solved." But TechCrunch reports that Nvidia's accounting draws a boundary around the data center itself, leaving out water consumed during electricity generation and chip manufacturing.

That omission is significant. Natural gas plants use 1.17 liters of water per kilowatt-hour generated; coal plants use 2.2 liters. Fossil fuels collectively supply roughly half of all data center power today, and the IEA projects they will account for more than 40% of new electricity needed to meet data center demand through 2030. By contrast, wind and solar use approximately 0.01 and 0.03 liters per kilowatt-hour, respectively. Nvidia's system may address one-quarter to one-third of a facility's total water footprint, according to TechCrunch's analysis, leaving the broader question tied directly to how data centers are powered.

SECTION 08

Creative AI

A24 and Google DeepMind Launch \$75M Multi-Year AI Filmmaking Partnership

Read the article: [Hollywood Reporter](#)

A24 and Google DeepMind have announced a multi-year research partnership in which the two companies will co-develop AI tools for use by A24's filmmakers, with those tools also feeding back into Google's broader ecosystem. DeepMind will invest \$75 million in the project. The deal is non-exclusive, leaving both parties free to collaborate with other studios and AI companies. Scott Belsky, hired by A24 in early 2025 to lead digital initiatives, will run the program with a team of roughly two dozen. The studio is already prototyping a storyboard tool under the arrangement.

The partnership marks DeepMind's first known collaboration with a full studio, having previously worked only with individual filmmakers such as Darren Aronofsky. For legal and creative industries watching AI's integration into media production, the deal is notable on several fronts: it advances a trend toward bespoke, brand-specific AI models rather than general-purpose generators, and it arrives amid vocal skepticism from prominent filmmakers about AI's role in creative work. Competitors including Netflix, Amazon Studios, and Lionsgate have already pursued similar arrangements.

SECTION 09

Higher Education

The Future of University Honor Codes in the AI Era

Read the article: [Inside Higher Ed](#)

Stanford and Princeton made headlines earlier this year when both universities approved the addition of proctored in-person exams, ending long-standing traditions of student-enforced honor systems. Stanford's change takes effect in September; Princeton's faculty approved a parallel policy in May. Both decisions were shaped in part by students' growing use of artificial intelligence, sanctioned and otherwise.

The shift reflects a broader crisis for honor codes at American universities. A 2025 Inside Higher Ed survey of more than 1,000 students found that 85 percent had used generative AI to complete coursework, with 19 percent reporting they used it to write essays or free responses. Academic integrity experts note that AI complicates traditional honor code enforcement because the social norms that discourage cheating off a classmate's exam carry little weight when a student is alone with a chatbot.

University of Virginia's student honor committee chair reports her institution does not use AI detection tools, citing unreliable results, and instead compares submitted work against a student's prior writing. Experts interviewed in the article disagree on whether adding faculty surveillance is sufficient, with some arguing that assessment methods themselves must change before honor codes can meaningfully adapt to the AI era.



This newsletter contains AI-generated content that may contain inaccuracies. For research or citation purposes, please read and cite the original source article, not this newsletter.