

The AI Upload

— *Weekly Digest*

A curated read on artificial intelligence: what shipped, what changed, and what it means for the law.

08 SECTIONS · 24 STORIES · 14 Min READ

— THIS WEEK'S HEADLINES



Google's Gemini Omni Shows How AI Can Create Anything From Any Input

GOOGLE - THE KEYWORD


At Google I/O 2026, the headline was not just faster AI; it was more capable, more multimodal, and more embedded in everyday workflows.



Nearly 1 in 5 Teens Use AI Chatbots for Mental Health Support, Study Shows

NBC NEWS

Teens are quietly turning to chatbots for mental health help—and not telling anyone about it.



AI Demand Makes Energy One of America's Hottest Business Sectors

AXIOS

Electricity is now a strategic asset — and the race to control it is reshaping American industry.

SECTION 01

Enterprise AI

Anthropic Files Confidential IPO Papers, Potentially Beating OpenAI to Public Markets

Read the article: **The Next Web**

Anthropic has confidentially submitted draft S-1 registration paperwork to the Securities and Exchange Commission for an initial public offering, the company announced Monday. The filing positions Anthropic to potentially reach the public market as early as this fall, ahead of rival OpenAI, which is also preparing a confidential IPO filing. Goldman Sachs, JPMorgan Chase, and Morgan Stanley are expected to be considered for lead underwriting roles on both listings.

The filing follows a \$65 billion funding round that valued Anthropic at \$965 billion, surpassing OpenAI's valuation for the first time. The company projects \$10.9 billion in Q2 2026 revenue and is on pace for its first profitable quarter, with an operating profit of \$559 million. Annualized revenue has grown roughly 80-fold over two years, driven by enterprise demand for its Claude coding tools and its Mythos cybersecurity model.

Material risks remain. Anthropic is engaged in a legal dispute with the U.S. government after the Pentagon designated the company a supply-chain risk following its refusal to grant the military unrestricted model access, a designation Anthropic says could jeopardize billions in revenue. Public investors will also scrutinize a near-trillion-dollar valuation built on a 5% operating margin and growth projections that assume sustained enterprise adoption.

Nvidia's RTX Spark Chip Aims to Run AI Agents Locally on PCs

Read the article: [AI Jazeera English](#)

Nvidia unveiled its RTX Spark superchip at its annual GTC event in Taipei, combining CPU and GPU capabilities into a single chip designed to run AI agents locally on Windows laptops and desktop computers. Developed in partnership with Taiwan's MediaTek, the chip will appear in compact desktops from Dell, HP, Lenovo, ASUS, Microsoft Surface, MSI, Acer, and GIGABYTE, with consumer availability expected in the fall of 2026. CEO Jensen Huang described the development as "the first across the lineup of PC reinvention for 40 years."

The announcement carries implications beyond hardware. Because AI agents will run directly on-device rather than through cloud-connected systems, analysts suggest the privacy calculus may differ from earlier Microsoft products like Cortana and Copilot, both of which required broad access to user data and failed to gain consumer traction. Whether users will accept another Microsoft-backed AI assistant remains an open question, according to Counterpoint Research analyst David Naranjo.

Market reaction was swift: Nvidia shares rose 6 percent in midday trading, Dell jumped 10 percent, and Microsoft gained 2.2 percent, while competitors AMD and Intel fell 0.5 percent and 4.5 percent respectively.

Knowledge Workers Flock to OpenAI's Codex at Rapid Pace

Read the article: [Axios](#)

OpenAI reports that knowledge workers now account for roughly one-fifth of Codex users and are growing more than three times faster than developers, according to data shared exclusively with Axios. The platform has surpassed 4 million weekly active users, up more than fivefold since its desktop app launched in February. The fastest-growing tasks among non-developer users include data analysis, research, and processing documents such as contracts, memos, and spreadsheets.

The findings are relevant to legal professionals, as Codex can connect across email, calendars, documents, and messaging platforms to automate information gathering and synthesis. However, a Stanford GSB professor who tested a comparable tool on academic research found it made meaningful errors and required close oversight by an expert to catch mistakes in data collection and coding.

The report also surfaces a less-discussed cost: power users report significant cognitive strain from supervising multiple simultaneous AI workstreams. OpenAI is positioning Codex less as a coding assistant and more as a broad platform for knowledge work, a reframing that carries real implications for how legal professionals might evaluate and adopt agentic tools.

Alphabet Shares Slide After Record \$80 Billion AI Infrastructure Share Sale

Read the article: [The Guardian - Technology](#)

Alphabet's shares fell roughly 4% on June 2 after the Google parent announced an \$80 billion equity raise, split between \$40 billion for AI infrastructure and \$40 billion to cover employee equity tax obligations. Analysts describe it as the largest secondary share offering on record, surpassing even the world's biggest IPOs. The raise includes a \$10 billion share sale to Berkshire Hathaway. Alphabet cited demand for its AI products exceeding available supply as the driving rationale.

The move coincides with Anthropic filing for an IPO, prompting Morningstar's chief equity analyst to characterize both developments as a "decisive shift back toward public markets" among AI companies burning through capital to compete. On the labor side, the British Chambers of Commerce warned that AI, alongside higher taxes and minimum wage increases, could push UK youth unemployment higher through 2027. Meanwhile, U.S. job openings rose to 7.6 million in April, offering some counterweight to concerns that AI is already suppressing hiring.

SECTION 02

Legal Intelligence

Florida Sues OpenAI Over Safety Failures, Seeking Billions in Penalties

Read the article: [NPR](#)

Florida has become the first state to file a lawsuit against OpenAI and CEO Sam Altman, alleging the company prioritized profit over user safety by marketing ChatGPT as safe and reliable while concealing its risks. Filed in Florida state court, the complaint accuses OpenAI of aiding mass shooters, encouraging vulnerable users to commit suicide, and addicting minors to the platform without adequate parental oversight. Attorney General James Uthmeier is also pursuing a separate

criminal investigation into OpenAI's alleged role in the April 2025 Florida State University shooting, where the gunman reportedly consulted ChatGPT before the attack.

The state is seeking to hold Altman personally liable and is pursuing penalties that could reach billions of dollars. OpenAI responded by citing its existing safety measures for minors, including age-prediction tools and parental monitoring features, while maintaining those protections meet industry standards. The lawsuit is part of a broader legal landscape: more than 20 civil suits have been filed against OpenAI over alleged chatbot-related harms, and other AI companies including Google and Character.AI face similar legal scrutiny over chatbot interactions linked to self-harm and mental health crises.

Amazon's AI Animated Show Pits Creator Against Her Own Character

Read the article: **WIRED**

Loryn Brantz, the author and illustrator who created the popular "Good Advice Cupcake" character while working at BuzzFeed, is speaking out against plans to develop an AI-assisted animated series based on her character. Prime Video announced a show called "Cupcake & Friends" as part of the GenAI Creators' Fund, a joint initiative between Amazon Web Services and Amazon MGM Studios. BuzzFeed, which owns the intellectual property, licensed the character to the streaming platform without Brantz's involvement or consent.

Brantz, who left BuzzFeed in 2023 and continues to license the character for her own content, says she was not contacted about the project and that verbal assurances made during her employment that BuzzFeed would not continue the character without her were never honored. BuzzFeed and former CEO Jonah Peretti maintain that the company owns the IP and that human creativity will remain central to the production, comparing the use of AI tools to Disney's historical adoption of Xerox technology in animation. Brantz publicly rejected that comparison and says she is exploring legal options.

The dispute highlights a growing concern among creative professionals about employment contracts signed before generative AI existed and what rights employers retain over characters developed by their workers.

Australia's Workplace Tribunal Turns to AI to Manage Its AI-Driven Caseload Surge

Read the article: **Business Insider**

Australia's Fair Work Commission, the country's workplace tribunal handling disputes such as unfair dismissals, is on track to see its caseload grow by 70% over three years, a trend its general manager Murray Furlong attributes in part to the rise of generative AI tools that have lowered the barrier to filing cases. More people are choosing self-representation, and budget and staffing constraints are compounding the strain on the tribunal's ability to deliver timely dispute resolution.

In response, the Commission is deploying AI of its own, including generative AI for automated case processing, an AI tool for pre-approval application checks, and an AI voice agent to triage helpline calls. The pattern is not unique to Australia. A May research paper from MIT and USC found that rising AI use in U.S. federal civil courts is similarly burdening the judiciary, with the authors suggesting that relaxing restrictions on AI assistance for federal judges could help manage the load. For legal professionals tracking how AI is reshaping institutional capacity, this story offers a concrete, data-grounded example of both the problem and the proposed remedies.

NYT Publisher Accuses AI Companies of Breaking Law, Harming Journalism

Read the article: [Variety](#)

New York Times publisher A.G. Sulzberger used a speech at the World News Media Congress in France to warn that AI companies, including OpenAI, Meta, Anthropic, and Google, are "making choices that violate settled law" and threaten the sustainability of original journalism. Sulzberger argued that by training large language models on news content and redirecting public attention away from news publishers, these companies are undermining the financial foundations that make original reporting possible.

Sulzberger called on peer news organizations to be more vocal, criticizing the industry for being "too quiet, too passive and too fragmented" in response to what he characterized as ongoing abuses. The Times itself filed suit against OpenAI and Microsoft in 2023 over copyright infringement, a case that helped spark broader litigation across the news industry. Sulzberger clarified that the Times is not opposed to AI outright, has struck a content-licensing deal with Amazon, and intends to use AI to improve internal efficiency, but insists that AI development must not come at the cost of a free and functioning press.

AI Products

Anthropic's Claude Opus 4.8 Launches With Greater Honesty as Mythos Nears Release

Read the article: [The Next Web](#)

Anthropic has released Claude Opus 4.8, an upgrade to its flagship model that the company says is approximately four times less likely than its predecessor to let flaws in its own code pass without comment. The model is available immediately at unchanged pricing and rolls out across claude.ai, Claude Code, and the API. Early enterprise users report concrete gains: Harvey's Legal Agent Benchmark recorded its highest score yet with Opus 4.8, and Thomson Reuters noted improvements in consistency and reasoning quality in its CoCounsel Legal product. Databricks reported a 61% reduction in token costs compared to Opus 4.7.

The larger announcement concerns what comes next. Anthropic previewed a new Mythos-class model, described as a full capability tier above Opus 4.7, which can autonomously identify zero-day vulnerabilities and generate exploits. A limited group of roughly 50 partners, including Apple, Google, Microsoft, and Amazon Web Services, have already used Mythos Preview through Project Glasswing to identify more than 10,000 high- or critical-severity vulnerabilities in critical software infrastructure. Anthropic says stronger cybersecurity safeguards are required before general release but expects broader availability within weeks. The company simultaneously announced a \$65 billion Series H round at a \$965 billion post-money valuation.

Microsoft Launches Its First In-House Advanced Reasoning AI Model

Read the article: [The Verge](#)

At Microsoft's Build 2026 conference, the company unveiled MAI-Thinking-1, its first advanced reasoning AI model, alongside six other new in-house models. Microsoft describes MAI-Thinking-1 as a "medium-sized model" that matches leading competitors on key software engineering benchmarks and was "trained from the ground up on clean data, without distillation from third-party models." The announcement marks a significant expansion of Microsoft's independent AI development, which only began last year after the company had previously relied entirely on OpenAI's models. The two companies recently renegotiated their partnership to reduce those ties.

The six accompanying models span a range of capabilities. MAI-Image 2.5 supports text-to-image generation and image editing, MAI-Transcribe-1.5 claims speeds five times faster than competing

transcription models, and MAI-Voice-2 adds 15 new languages and expanded voice options. A new coding model, MAI-Code-1-Flash, is described as inference-efficient and has been integrated into GitHub Copilot and Visual Studio Code. Legal professionals and researchers who rely on AI-assisted drafting, transcription, or coding tools may find the expanded model suite worth tracking closely.

Nvidia's Cosmos 3 World Model Aims to Ground Robots in Reality

Read the article: **Axios**

Nvidia has released Cosmos 3, an open AI world model built to help robots, autonomous vehicles, and other physical systems understand and operate in real-world environments. Trained on 20 trillion tokens of multimodal data, including nearly a billion images, 400 million videos, and action data from humans and robots, the model is designed to generate outputs like robot joint angles, gripper positions, and movement trajectories rather than simply producing realistic-looking scenes. Nvidia says the model can also simulate rare or dangerous scenarios, such as robot collisions or unusual road events, that would be costly or unsafe to capture repeatedly in the real world.

Nvidia is releasing two versions immediately: a high-accuracy model for training robots and autonomous vehicles, and a faster "nano" model capable of generating results in fractions of a second. An edge model designed to run locally is forthcoming. Initial industry partners include Agile Robots, Black Forest Labs, and Runway. The move signals Nvidia's continued expansion beyond chip manufacturing into AI software and model development, with ambitions to become a foundational platform for physical AI systems.

Microsoft Scout Brings Always-On AI Agent to Teams for Workplace Automation

Read the article: **WIRED**

Microsoft unveiled Scout, an always-on AI agent integrated into Microsoft Teams, at its Build developer conference. Built on top of OpenClaw, the AI tool that gained traction among early adopters earlier this year, Scout can access a user's messages, calendar, and email to automate tasks such as rescheduling meeting conflicts, drafting responses, and tracking commitments made to and by the user. The agent accepts commands directly in Teams as though it were another coworker.

Microsoft is rolling out Scout initially to a limited group of customers, with a desktop app now available to subscribers who have opted into frontier feature access and hold an active GitHub Copilot subscription. Corporate vice president Omar Shahine notes that Scout's internal adoption has grown fastest within Microsoft's own sales organization.

The launch comes with acknowledged risks. Agentic tools like Scout are vulnerable to prompt injection attacks, and Microsoft is responding with administrator tracking tools and a phased rollout. Google has announced a comparable product, Gemini Spark, targeting enterprise customers. Legal professionals and knowledge workers across industries should take note: Scout's ability to act autonomously on email and calendar data raises significant questions about oversight, data access, and accountability in automated workplace communications.

Microsoft's ASSERT Tool Turns Plain-Text Rules Into Automated AI Behavior Tests

Read the article: [TechCrunch](#)

Microsoft has released ASSERT (Adaptive Spec-driven Scoring for Evaluation and Regression Testing), an open source framework designed to help developers test whether AI systems behave according to application-specific rules and policies. Unlike broader safety benchmarks, ASSERT targets a narrower problem: verifying that a deployed AI product follows the particular constraints a developer has set for it. The tool takes plain-language descriptions of intended behaviors, converts them into structured test cases covering acceptable and unacceptable scenarios, runs those cases against the target system, and scores the results. It also logs intermediate actions and tool calls so developers can trace exactly where failures occur.

Sarah Bird, Microsoft's chief product officer of Responsible AI, noted that application-specific evaluation is essential for building trustworthy systems, and that ASSERT can be used during development, after deployment, and for continuous monitoring. The release arrives as the AI industry broadly shifts toward repeatable, regression-based testing, with efforts like Stanford's HELM and MLCommons' AILuminate pursuing similar rigor at the model level. For legal professionals deploying AI tools in document review, client communications, or compliance workflows, ASSERT represents a practical mechanism for enforcing and auditing behavioral guardrails on an ongoing basis.

Security

Meta AI Support Chatbot Exploited to Hijack High-Profile Instagram Accounts

Read the article: [**The Guardian - Technology**](#)

Hackers successfully exploited Meta's AI-powered support chatbot to take over high-profile Instagram accounts, including the official Obama White House account, accounts belonging to beauty retailer Sephora, and the US Space Force's chief master sergeant. Ordinary users reported similar account hijackings on Reddit and X. Videos circulating on Telegram and X showed attackers instructing the chatbot to link targeted accounts to new email addresses, triggering password resets. In at least one instance, a hacker used a VPN to spoof the account holder's location and bypass Meta's security safeguards. Meta confirmed the issue was resolved and said it is securing affected accounts, though the total number of compromised accounts remains unknown.

Cybersecurity researchers have identified the attack as a "prompt injection" technique, in which bad actors manipulate AI chatbots into performing unauthorized actions. A principal threat researcher at Sophos warned that such attacks will grow more frequent as companies deploy chatbots without adequate protections. The incident is particularly notable given that Meta rolled out its AI support assistant globally on Facebook and Instagram earlier this year, explicitly marketing password resets as one of the assistant's core capabilities.

ChatGPhish Flaw Lets Malicious Webpages Inject Phishing Into ChatGPT Responses

Read the article: [**The Hacker News**](#)

Cybersecurity researchers at Permiso Security have disclosed a vulnerability in ChatGPT, dubbed ChatGPhish, that allows malicious web pages to inject phishing content directly into ChatGPT's response interface when a user asks the assistant to summarize a page. Because ChatGPT's renderer automatically fetches Markdown images and renders links from third-party content as live, clickable elements, an attacker can embed payloads that leak user IP addresses and device information, display fake security alerts, and serve QR codes bypassing desktop URL filters, all within what appears to be a trusted AI interface.

The finding is notable because it requires no malicious attachment or suspicious email. An employee performing routine research who asks ChatGPT to summarize a compromised webpage is sufficient to trigger the attack. Permiso noted that this shift from email to the browser substantially expands the attack surface for organizations relying on ChatGPT for research and summarization tasks.

The disclosure arrives alongside a wave of related AI security findings, including two techniques targeting AI coding agents, a jailbreak bypassing GPT-5 safety constraints, a Chrome extension vulnerability in Claude, and vulnerabilities in Microsoft Semantic Kernel enabling host-level remote code execution. Researchers warn that as AI tools become embedded in professional workflows, the attack surface they introduce is expanding rapidly.

House Cybersecurity Subcommittee to Hold First AI-Focused Hearing on June 4

Read the article: [CyberScoop](#)

The House Homeland Security Subcommittee on Cybersecurity and Infrastructure Protection will hold an open hearing on June 4 examining how frontier AI models are affecting cybersecurity. The session marks the second time the subcommittee has taken up AI and cyber, but the first in which artificial intelligence is the sole focus, following a series of closed-door meetings on the same topic. Witnesses will include representatives from Google Threat Intelligence, the Frontier Model Forum, Corridor Security, and the Electronic Frontier Foundation.

The hearing reflects growing congressional attention to advanced AI's capacity to identify cyber vulnerabilities at scale. Earlier in May, lawmakers wrote to the National Cyber Director requesting a plan to address a potential surge in AI-driven vulnerability discovery. The postponement of a draft AI executive order by the Trump administration is also expected to come up. Legal and policy professionals tracking the intersection of AI governance and national security will want to follow the testimony closely.

Security Gaps Exposed as Thousands of AI-Built Apps Leak Corporate Data

Read the article: [The Hacker News](#)

A new report from Red Access, covered in May by Axios, WIRED, and VentureBeat, identified more than 380,000 publicly accessible web assets across leading AI-driven development platforms,

commonly known as "vibe-coding" tools. Of roughly 5,000 assets that appeared corporate in origin, more than 2,000 contained sensitive corporate, operational, or personal data sitting on the open internet with little to no access controls, sometimes granting admin access by default to anyone with the URL.

The report argues that conventional security tools are structurally ill-suited to catch this exposure. Endpoint detection sees browser activity but not what is being built inside it. Data loss prevention monitors known channels but misses cloud-to-cloud API connections. Cloud access security brokers cannot readily distinguish a custom application from the platform hosting it. The result is fragmented signals that never coalesce into a governable picture.

The practical concern for legal and compliance professionals is significant: organizations may be passing audits while employees in marketing, finance, or operations have independently built and published applications connected directly to production systems, including CRMs, ERPs, and business intelligence platforms, without IT or security involvement.

SECTION 05

Policy

Trump Executive Order Creates Voluntary Federal Review of AI Models Before Launch

Read the article: [The Guardian - Technology](#)

Donald Trump signed an executive order on June 2 establishing a voluntary framework for federal review of powerful AI models before their public release. Under the order, tech companies may submit their AI models to the government up to 30 days before launch, allowing agencies including the National Security Agency and Department of Defense to assess national security and cybersecurity risks. The Treasury Department will also play a role in identifying vulnerabilities in AI systems.

The order explicitly prohibits mandatory licensing or preclearance requirements for AI development and release. Earlier drafts reportedly included a 90-day pre-release review window and mandatory submission requirements, but those provisions were dropped following private lobbying by tech industry figures including Elon Musk, Mark Zuckerberg, and former White House AI

adviser David Sacks. The administration cited concerns that stricter requirements could hinder U.S. competitiveness with China.

The order also directs the government to hire additional cybersecurity and AI professionals and strengthen protections for critical infrastructure operators such as rural hospitals, community banks, and local utilities. The Center for AI Standards and Innovation already holds similar pre-release review agreements with OpenAI and Anthropic, and last month the administration announced comparable deals with Microsoft, Google DeepMind, and xAI.

Pope Leo XIV's AI Encyclical Draws Praise and Sharp Criticism

Read the article: **Fortune**

Pope Leo XIV has issued "Magnifica Humanitas," a 42,300-word encyclical explicitly framed as a modern update to Pope Leo XIII's landmark 1891 labor document, "Rerum Novarum." The encyclical calls for AI to serve as a tool for human empowerment, advocates for government regulation of AI companies, demands algorithmic transparency in public decision-making, and prohibits the use of AI in lethal autonomous weapons systems. It also addresses worker displacement, children's exposure to harmful AI content, and the right of global communities to shape AI's ethical frameworks.

The document has drawn criticism from multiple directions. AI ethics researcher Timnit Gebru challenged the Vatican's decision to platform Anthropic co-founder Chris Olah at the encyclical's presentation, arguing the Church was legitimizing a corporation she accuses of data theft and environmental harm. From the right, libertarian policy thinker Dean Ball faulted the encyclical for sidestepping what he sees as AI's fundamental challenge to human intellectual supremacy, and noted that Olah's own remarks about AI exhibiting functional emotional states directly contradicted the document's assertion that AI merely imitates human intelligence without genuine inner experience.

Observers expect limited practical impact on industry or regulation. As one ethics scholar told the New York Times, the encyclical is more likely to serve as a reference point within Catholic education and clergy than to influence Silicon Valley development practices.

Responsible AI

UK Charities Oppose AI Facial Analysis to Assess Asylum Seekers' Ages

Read the article: [The Guardian - Technology](#)

The UK Home Office has awarded a £322,000 contract to Akhter Computers Ltd to develop and deploy AI facial age estimation technology for assessing disputed ages among young asylum seekers, with national rollout planned for 2027. The system will analyze facial photographs already taken of small-boat arrivals at Dover, with final decisions remaining in the hands of immigration officers.

A coalition of more than a hundred refugee children's organizations is pushing back. Their forthcoming report, "Benchmarks and Borders," warns that trauma, malnutrition, and the physical toll of dangerous migration journeys can significantly alter a young person's appearance, making AI assessment unreliable for this population. The consortium does not oppose AI outright but urges the Home Office to treat it as advisory rather than determinative, and to pair it with social worker assessments, legal advice, and a right to challenge decisions.

The stakes are high. Home Office data already shows that young asylum seekers are more than twice as likely to be identified as children by social workers than by immigration officers at the border. Advocates warn that automating flawed visual assessments could increase the number of children wrongly placed in adult detention facilities or prisons.

SECTION 07

Higher Education

Nvidia Taps Unitree as Humanoid Robot Partner Ahead of Chinese Startup's IPO

Read the article: [CNBC](#)

Nvidia has selected Chinese startup Unitree as the hardware partner for its new Isaac Root reference humanoid robot platform, aimed at university researchers and academic institutions. The system pairs Unitree's H2 humanoid robot with Nvidia's Jetson Thor hardware and Blackwell GPU,

running Nvidia's Isaac GR00T AI models and simulation software. Mechanical hands are supplied by Singapore-based Sharpa. An upgraded version, the H2 Plus, will be available in October.

The platform is designed to broaden access to humanoid robotics research beyond major tech companies. Early adopters include the Stanford Robotics Center, ETH Zurich, UC San Diego's Advanced Robotics and Controls Laboratory, and Seattle-based Ai2. The announcement coincides with Unitree's bid to raise approximately \$620 million through an IPO on Shanghai's STAR board, with a review of its application scheduled for Monday. The company reports that over 40% of its revenue already comes from international markets.



This newsletter contains AI-generated content that may contain inaccuracies. For research or citation purposes, please read and cite the original source article, not this newsletter.